

Summary

GNNs can be efficiently employed to solve bandit problems on a large domain of large graphs, e.g. for drug discovery. They save you from exponential dependency on the number of nodes.

- Problems such as protein design, molecule and drug discovery, involve solving

$$G^* \in \arg \max_{G \in \mathcal{G}} f^*(G)$$

where f^* is unknown, and sampling from it is costly.

- Challenges in such applications are scaling to large domains, and to graphs with many nodes.
- We model this as bandit optimization on graphs, and propose how GNNs can be used to design scalable algorithms.

Problem Setting

- Iteratively evaluate the function via noisy observations

$$y_t = f^*(G_t) + \epsilon_t$$

i.i.d. zero-mean sub-Gaussian noise

- \mathcal{G} Finite set of undirected graphs with N nodes
Each graph has node features $\mathbf{h}_G = (\mathbf{h}_{G,j})_{j=1}^N \in \mathbb{R}^{Nd}$

- f^* Real valued, regular, invariant to node permutations

$$f^*(c \cdot G) = f^*(G) \quad \forall G \in \mathcal{G}, c \in P_N$$

$$\|f^*\|_{\bar{k}} \leq B$$

- Choose a kernel that is permutation invariant, but also efficient to compute

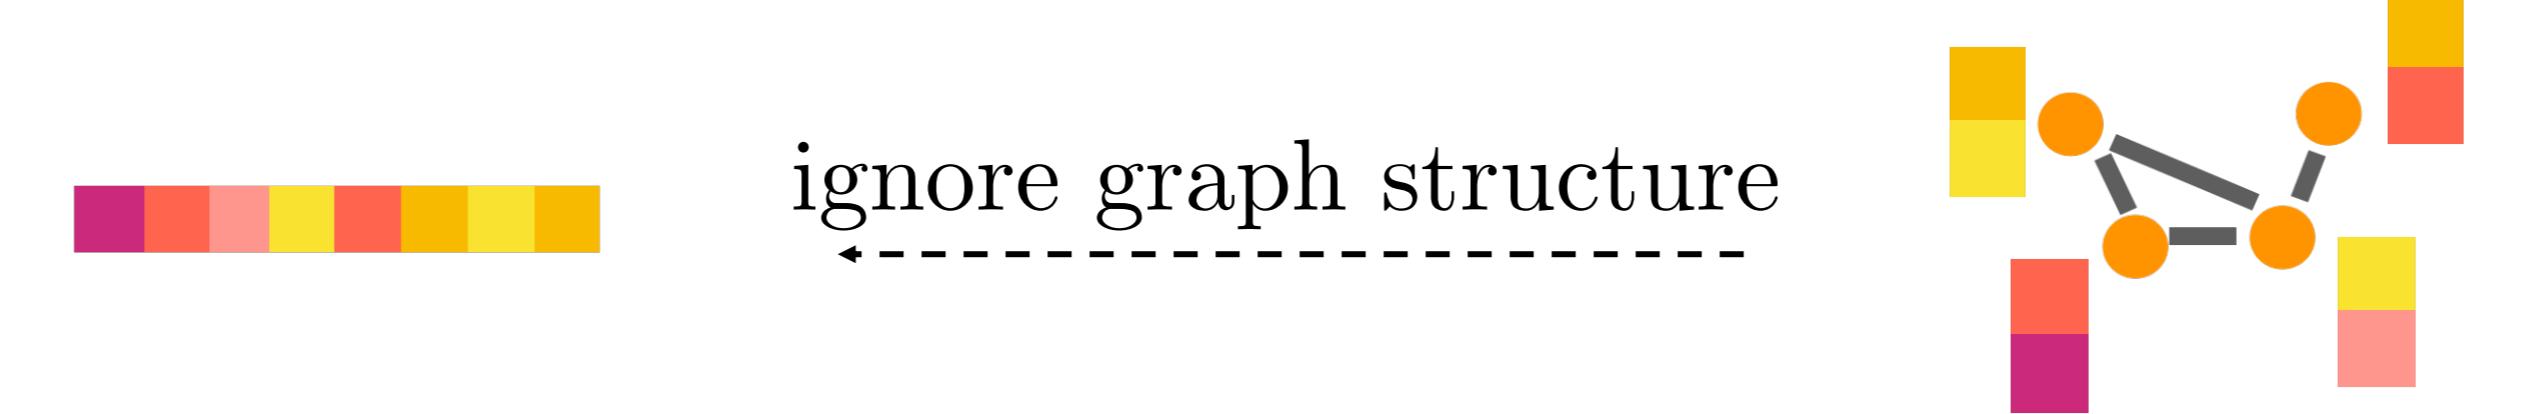
$$\bar{k}(G, G') = \frac{1}{|P_N|^2} \sum_{c, c' \in P_N} k(\bar{\mathbf{h}}_{c \cdot G}, \bar{\mathbf{h}}_{c' \cdot G'})$$

- Setting k as an NTK creates symmetries which gives

$$\bar{k} = k_{\text{GNN}} \quad k = \frac{1}{N} \sum_{j=1}^N k_{\text{NN}}(\bar{\mathbf{h}}_{G,j}, \bar{\mathbf{h}}_{G',j})$$

- The Graph Neural Tangent Kernel is expressive, efficient to compute, and additive permutation invariant.

Permutation Invariance Trick



Many block-permutations of this vector yield the same reward. If the estimator is agnostic to this structure,

$$\hat{f}(\text{permuted vector}) \neq \hat{f}(\text{original vector}) \Rightarrow \text{Sample inefficiency}$$

$$f^*(\text{permuted graph}) = f^*(\text{original graph})$$

→ Use a reward estimator which is invariant to these block permutations.

GNN with 1 conv layer, $f_{\text{GNN}}(G; \theta) = f_{\text{GNN}}(c \cdot G; \theta)$

GNN Confidence Sets

- Train $f_{\text{GNN}}(G; \theta)$ to estimate $f^*(G)$

Lazy initialization + SGD on

$$\mathcal{L}(\theta) = \frac{1}{t} \sum_{i < t} (f_{\text{GNN}}(G_i; \theta) - y_i)^2 + m\lambda \|\theta - \theta^0\|_2^2$$

$$\hat{\mu}_{t-1}(G) := f_{\text{GNN}}(G; \hat{\theta}_{t-1})$$

- Use $\nabla_{\theta} f_{\text{GNN}}$ to capture the uncertainty over the outputs

$$\hat{\sigma}_{t-1}^2(G) := \frac{\nabla f_{\text{GNN}}^T(G)}{\sqrt{m}} \left(\lambda \mathbf{I} + \mathbf{H}_{t-1} \right)^{-1} \frac{\nabla f_{\text{GNN}}(G)}{\sqrt{m}}$$

- Construct Confidence sets,

$$\mathcal{C}_{t-1}(G, \delta) = [\hat{\mu}_{t-1}(G) \pm \beta_t \hat{\sigma}_{t-1}(G)]$$

$$\beta_t \approx \sqrt{2B} + \frac{\sigma}{\sqrt{\lambda}} \sqrt{2 \log 2|\mathcal{G}|/\delta}$$

Theorem

GNN confidence sets are valid if the used network is wide enough, i.e. with probability greater than $1 - \delta$

$$f^*(G) \in \mathcal{C}_{t-1}(G, \delta), \quad \forall G \in \mathcal{G}$$

- Confidence sets as a proxy for the reward!

Regret Guarantee

- Cumulative regret as a measure of sample efficiency

$$R_T = \sum_{t=1}^T f^*(G^*) - f^*(G_t)$$

$$R_T/T \rightarrow 0 \text{ as } T \rightarrow \infty$$

- GNN-PE

- Has episodic structure
- Uses $\mathcal{C}_{t-1}(G, \delta)$
- To maintain set of plausible maximizer graphs
- To choose the next graph

Theorem

Suppose $f^* \in \mathcal{H}_{k_{\text{GNN}}}$ with a norm bounded by B .

Then with probability greater than $1 - \delta$,

$$R_T = \tilde{\mathcal{O}} \left(T^{\frac{2d-1}{2d}} \log^{\frac{1}{2d}} T \left(B + \sqrt{\log |\mathcal{G}|/\delta} \right) \right)$$

if the used GNN is wide enough.

$\log N$ dependency

- Naively using a NN gives

$$\tilde{\mathcal{O}} \left(T^{\frac{2Nd-1}{2Nd}} \log^{\frac{1}{2Nd}} T \right)$$

Experiments

- Training GNNs in lazy regime is challenging. Stopping criterion for gradient descent plays a crucial role in achieving sublinear regret.

- We devise a history-dependent stopping criterion.

