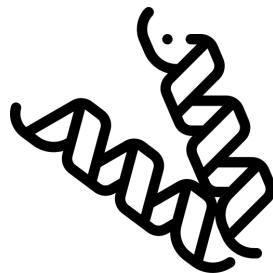


Graph Neural Network Bandits

Parnian Kassraie, Andreas Krause, Ilija Bogunovic



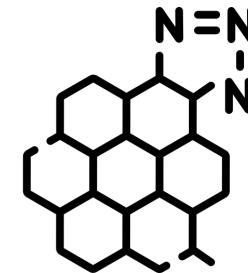
Learning on Graph Structured Data



Protein
Design

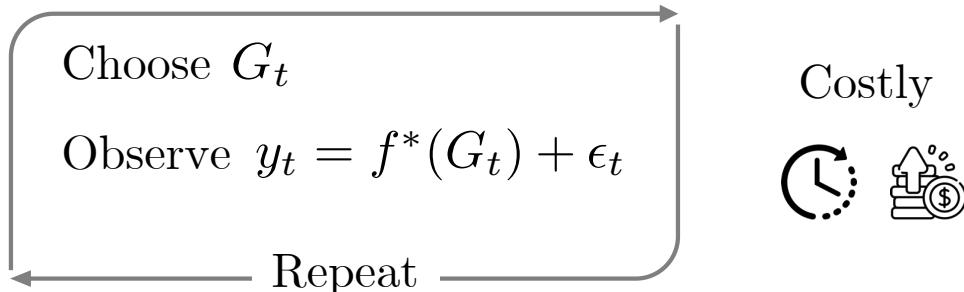


Drug
Discovery



Molecule
Synthesis

$$G^* \in \arg \max_{G \in \mathcal{G}} f^*(G)$$



Costly → Need to be sample efficient!
→ Model as a bandit problem


Problem Setting

\mathcal{G} Finite set of undirected graphs with N nodes
Each graph has node features

f^* Real-valued and regular (contained within an RKHS)
Invariant to node permutations

$$f^*(c \cdot G) = f^*(G)$$

if you indexed the graphs in your dataset in a different way, it would not matter.




Can you use GNNs to efficiently maximize such functions on such domains?

Bandit Objective

$$R_T = \sum_{t=1}^T f^*(G^*) - f^*(G_t)$$

Sublinear if converges to maxima
Small if sample efficient

Why a GNN?



Structure-agnostic reward estimator

while,

$$\hat{f}(\text{---}) \neq \hat{f}(\text{---})$$
$$f^*(\text{---}) = f^*(\text{---})$$

\Rightarrow Sample inefficiency



Use a reward estimator which is invariant to these block permutations,
GNN with 1 conv layer,

$$f_{\text{GNN}}(G; \boldsymbol{\theta}) = f_{\text{GNN}}(c \cdot G; \boldsymbol{\theta}).$$

How to use GNNs

1) Train $f_{\text{GNN}}(G; \boldsymbol{\theta})$ to estimate $f^*(G)$

$$\hat{\mu}_{t-1}(G) := f_{\text{GNN}}(G; \hat{\boldsymbol{\theta}}_{t-1})$$

run SGD on

$$\mathcal{L}(\boldsymbol{\theta}) = \frac{1}{t} \sum_{i < t} \left(f_{\text{GNN}}(G_i, \boldsymbol{\theta}) - y_i \right)^2 + m\lambda \|\boldsymbol{\theta} - \boldsymbol{\theta}^0\|_2^2$$

2) Use $\nabla_{\boldsymbol{\theta}} f_{\text{GNN}}$ to capture the uncertainty over these estimates

$$\hat{\sigma}_{t-1}^2(G) := \frac{\nabla f_{\text{GNN}}^T(G)}{\sqrt{m}} \left(\lambda \mathbf{I} + \mathbf{H}_{t-1} \right)^{-1} \frac{\nabla f_{\text{GNN}}(G)}{\sqrt{m}}$$

gram matrix \mathbf{H}_{t-1}

width m

3) Use confidence sets as a guide to choose actions

$$\mathcal{C}_{t-1}(G, \delta) = [\hat{\mu}_{t-1}(G) \pm \beta_t \hat{\sigma}_{t-1}(G)]$$

Because:

Theorem

GNN confidence sets are valid if the used network is wide enough, i.e. with high probability

$$f^*(G) \in \mathcal{C}_{t-1}(G, \delta), \quad \forall G \in \mathcal{G}$$

Our algorithm

- GNN-PE**
- ▶ Has episodic structure
 - ▶ Uses $\mathcal{C}_{t-1}(G, \delta)$
 - ▶ To maintain set of plausible maximizer graphs
 - ▶ To choose the next graph

Theorem

Suppose $f^* \in \mathcal{H}_{k_{\text{GNN}}}$ and has a bounded norm.
If the used GNN is wide enough, then with high probability

$$R_T = \tilde{\mathcal{O}} \left(T^{\frac{2d-1}{2d}} \log^{\frac{1}{2d}} T \right)$$

$\log(N)$ and $\sqrt{\log(|\mathcal{G}|)}$ dependency

Naively using Neural UCB gives


$$\tilde{\mathcal{O}} \left(T^{\frac{2Nd-1}{2Nd}} \log^{\frac{1}{2Nd}} T \right)$$

T: the bandit horizon,
N: the number of nodes in each graph,
d: the dimension of node features

Experiments


Domain: Erdos-Renyi random graphs

Objective: Sampled from $\text{GP}(0, k_{\text{GNN}})$



- ✓ Outperforms NN methods
- ✓ Scales well to domains of large graphs
- ✓ Scales well to large domains of graphs

Checkout the paper for more



Thank you.